

Transaction: a set of logical operations

$R(A)$


$A = A - 100$


$\rightarrow$   
 $W(A)$   $\times$

$R(B)$

$B = B + 100$

$W(B)$





ACID properties: 4 important properties of a transaction

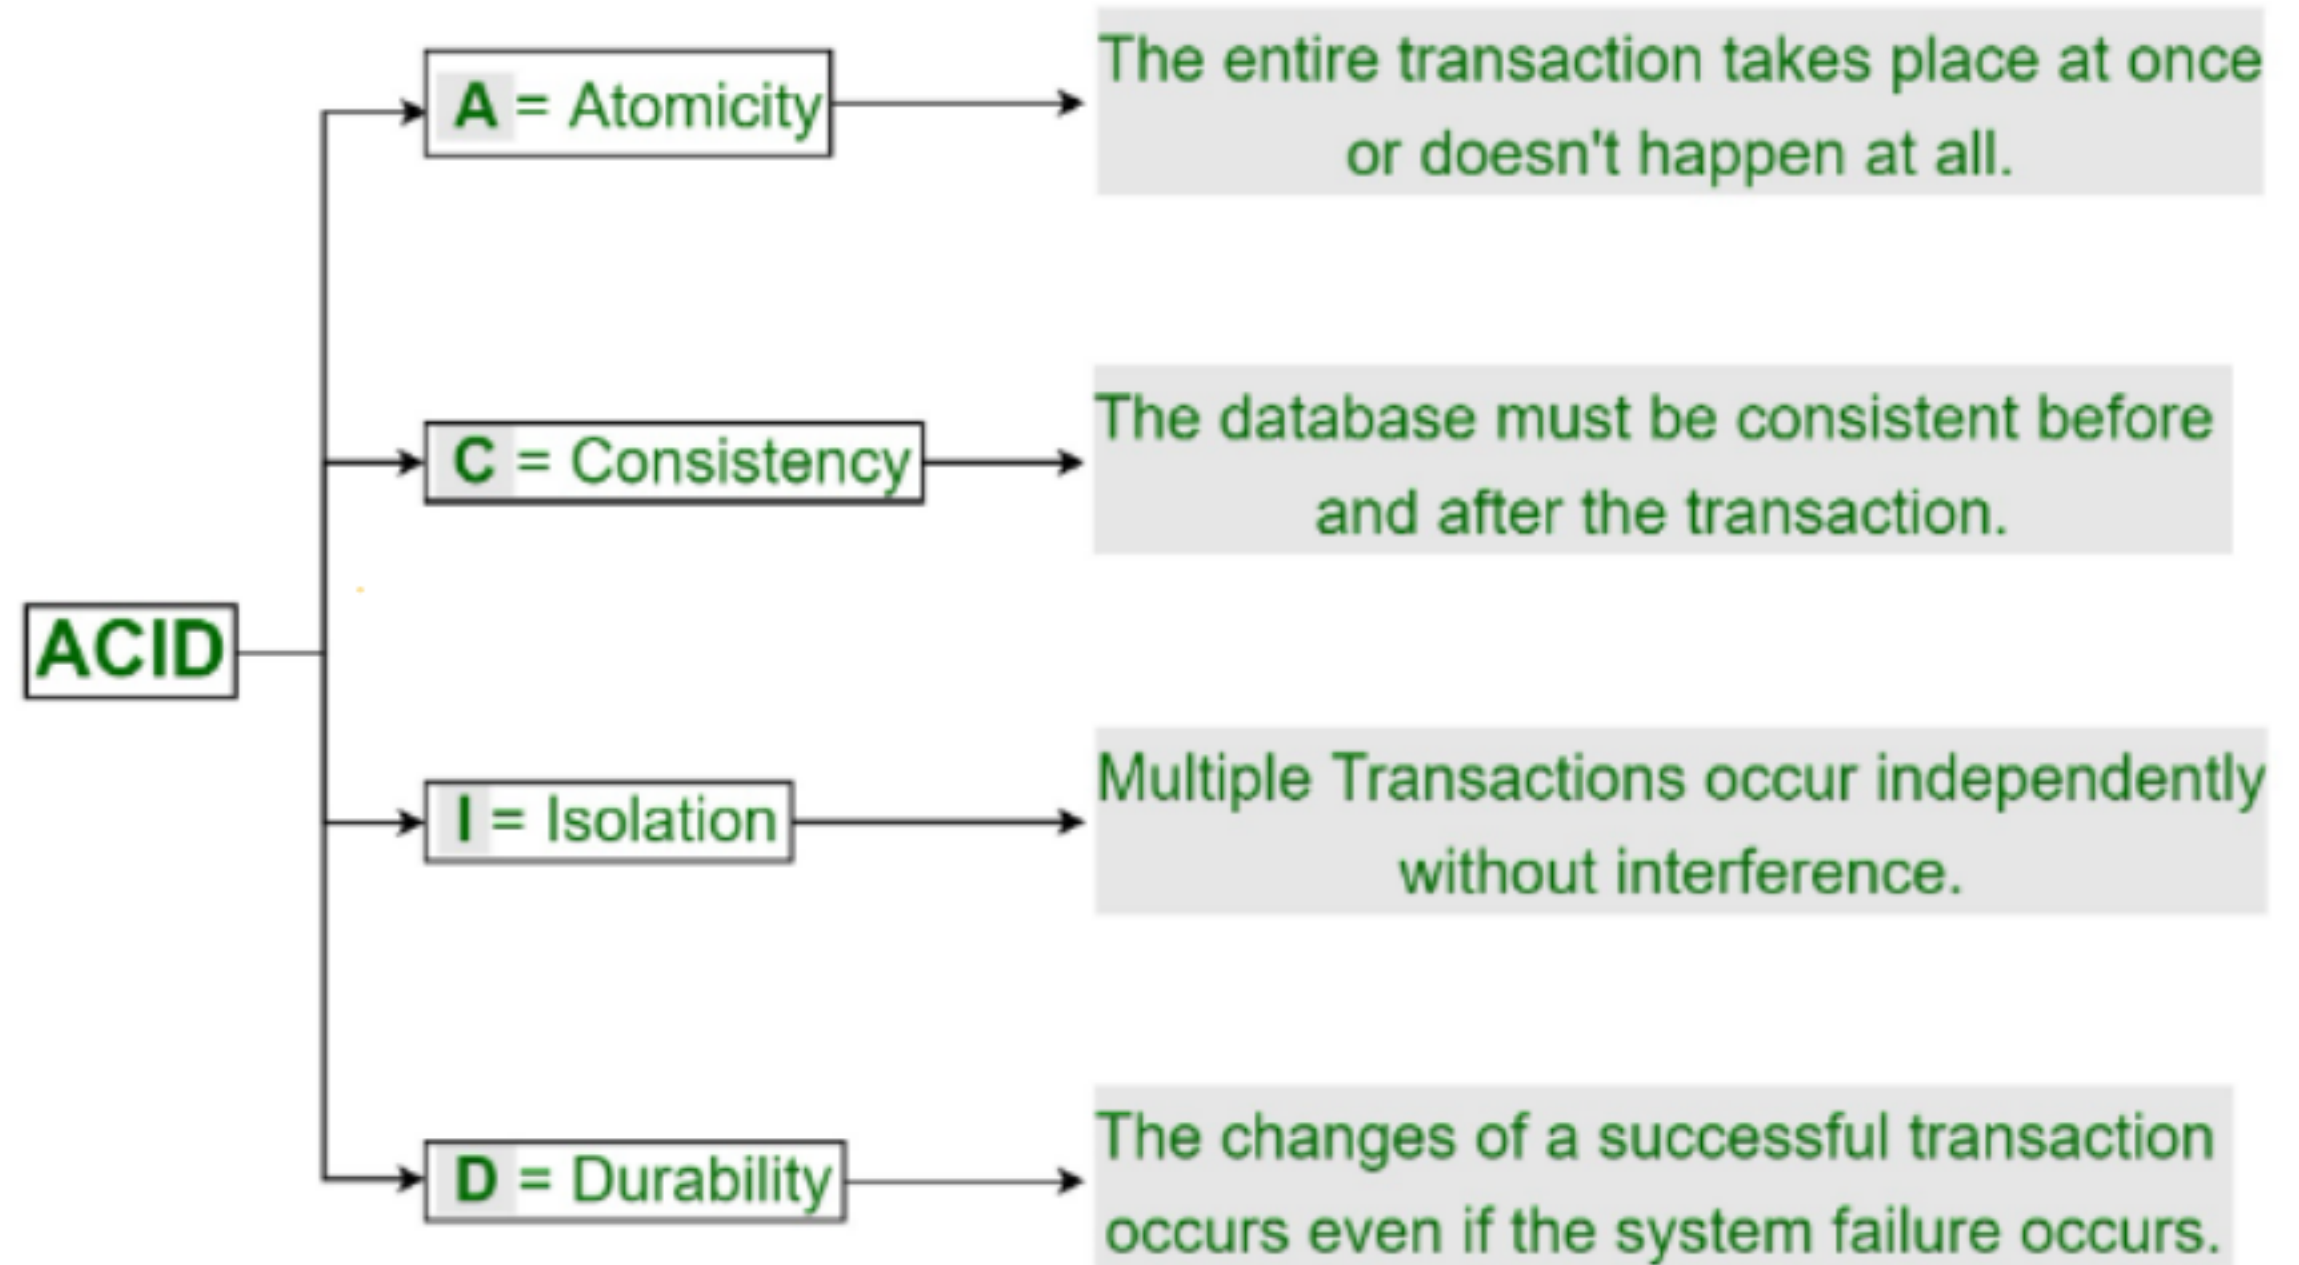
Atomicity

Consistency

Isolation

Durability

## ACID Properties in DBMS





- Atomicity ensures

- Either all the operations are executed successfully or none of them
- It involves two operations:
  - • Abort: If a transaction aborts, changes made to database are not visible
  - • Commit: If a transaction commits, changes made are visible

Example:

$T = \{R(A); A = A + 100; W(A);\}$

This transaction is about: credit \$100 to Account A.

If all 3 operations are executed successfully, commit operation needs to be performed

If T fails after Operation 2, rollback operation needs to be performed

- Concurrent execution of database transactions in a multi-user system environment means (more than 1 user can access the same database at the same time)

Consistency

- Consistency ensures
  - Concurrent transactions must not leave the database in an inconsistent state

Example:

$T_1 = \{R(A); A = A - 500; W(A);\}$   $T_2 = \{R(A); A = A + 1000; W(A);\}$

$T_1$  and  $T_2$  are two concurrent transactions.

Suppose Initially,  $A = 1000$ .

$T_1$  reads  $A(1000)$  and stores it in its local buffer. Then  $T_2$  reads  $A(1000)$  and also stores it in its local buffer.

$T_1$  performs  $A = A - 500$  and 500 is stored in local buffer of  $T_1$ . Then  $T_2$  performs  $A = A + 1000$  and 2000 is stored in buffer of  $T_2$ .

$T_1$  writes the value from its buffer back to database. Then  $T_2$  writes the value from its buffer back to database.

Finally,  $A = 2000$  in the database

It leads to an inconsistent state in a database



- • Isolation ensures
- Result of a transaction is not be visible to others before it is committed

Example:

→  $T_1 = \{R(A); A = A - 500; \underline{W(A)};\} \rightarrow \underline{500}$

→  $T_2 = \{R(A); A = A + 1000; \underline{W(A)};\} \underline{1500}$

Suppose Initially,  $A = 1000$ .

After successful execution of  $T_1$ ,  $A = 500$

After successful execution of  $T_2$ ,  $A = 1500$

⊗ → What happens  $T_1$  fails just before committing?

We need to rollback  $T_2$  as well, since  $T_2$  used the value of  $A$  generated by  $T_1$

⊗ → Conclusion: A transaction results should not made visible to other transactions before it commits



# Durability

- Durability ensures
  - Once database has committed a transaction, the changes made by the transaction are permanent

Example:

$T = \{R(A); A = A - 500; W(A);\}$

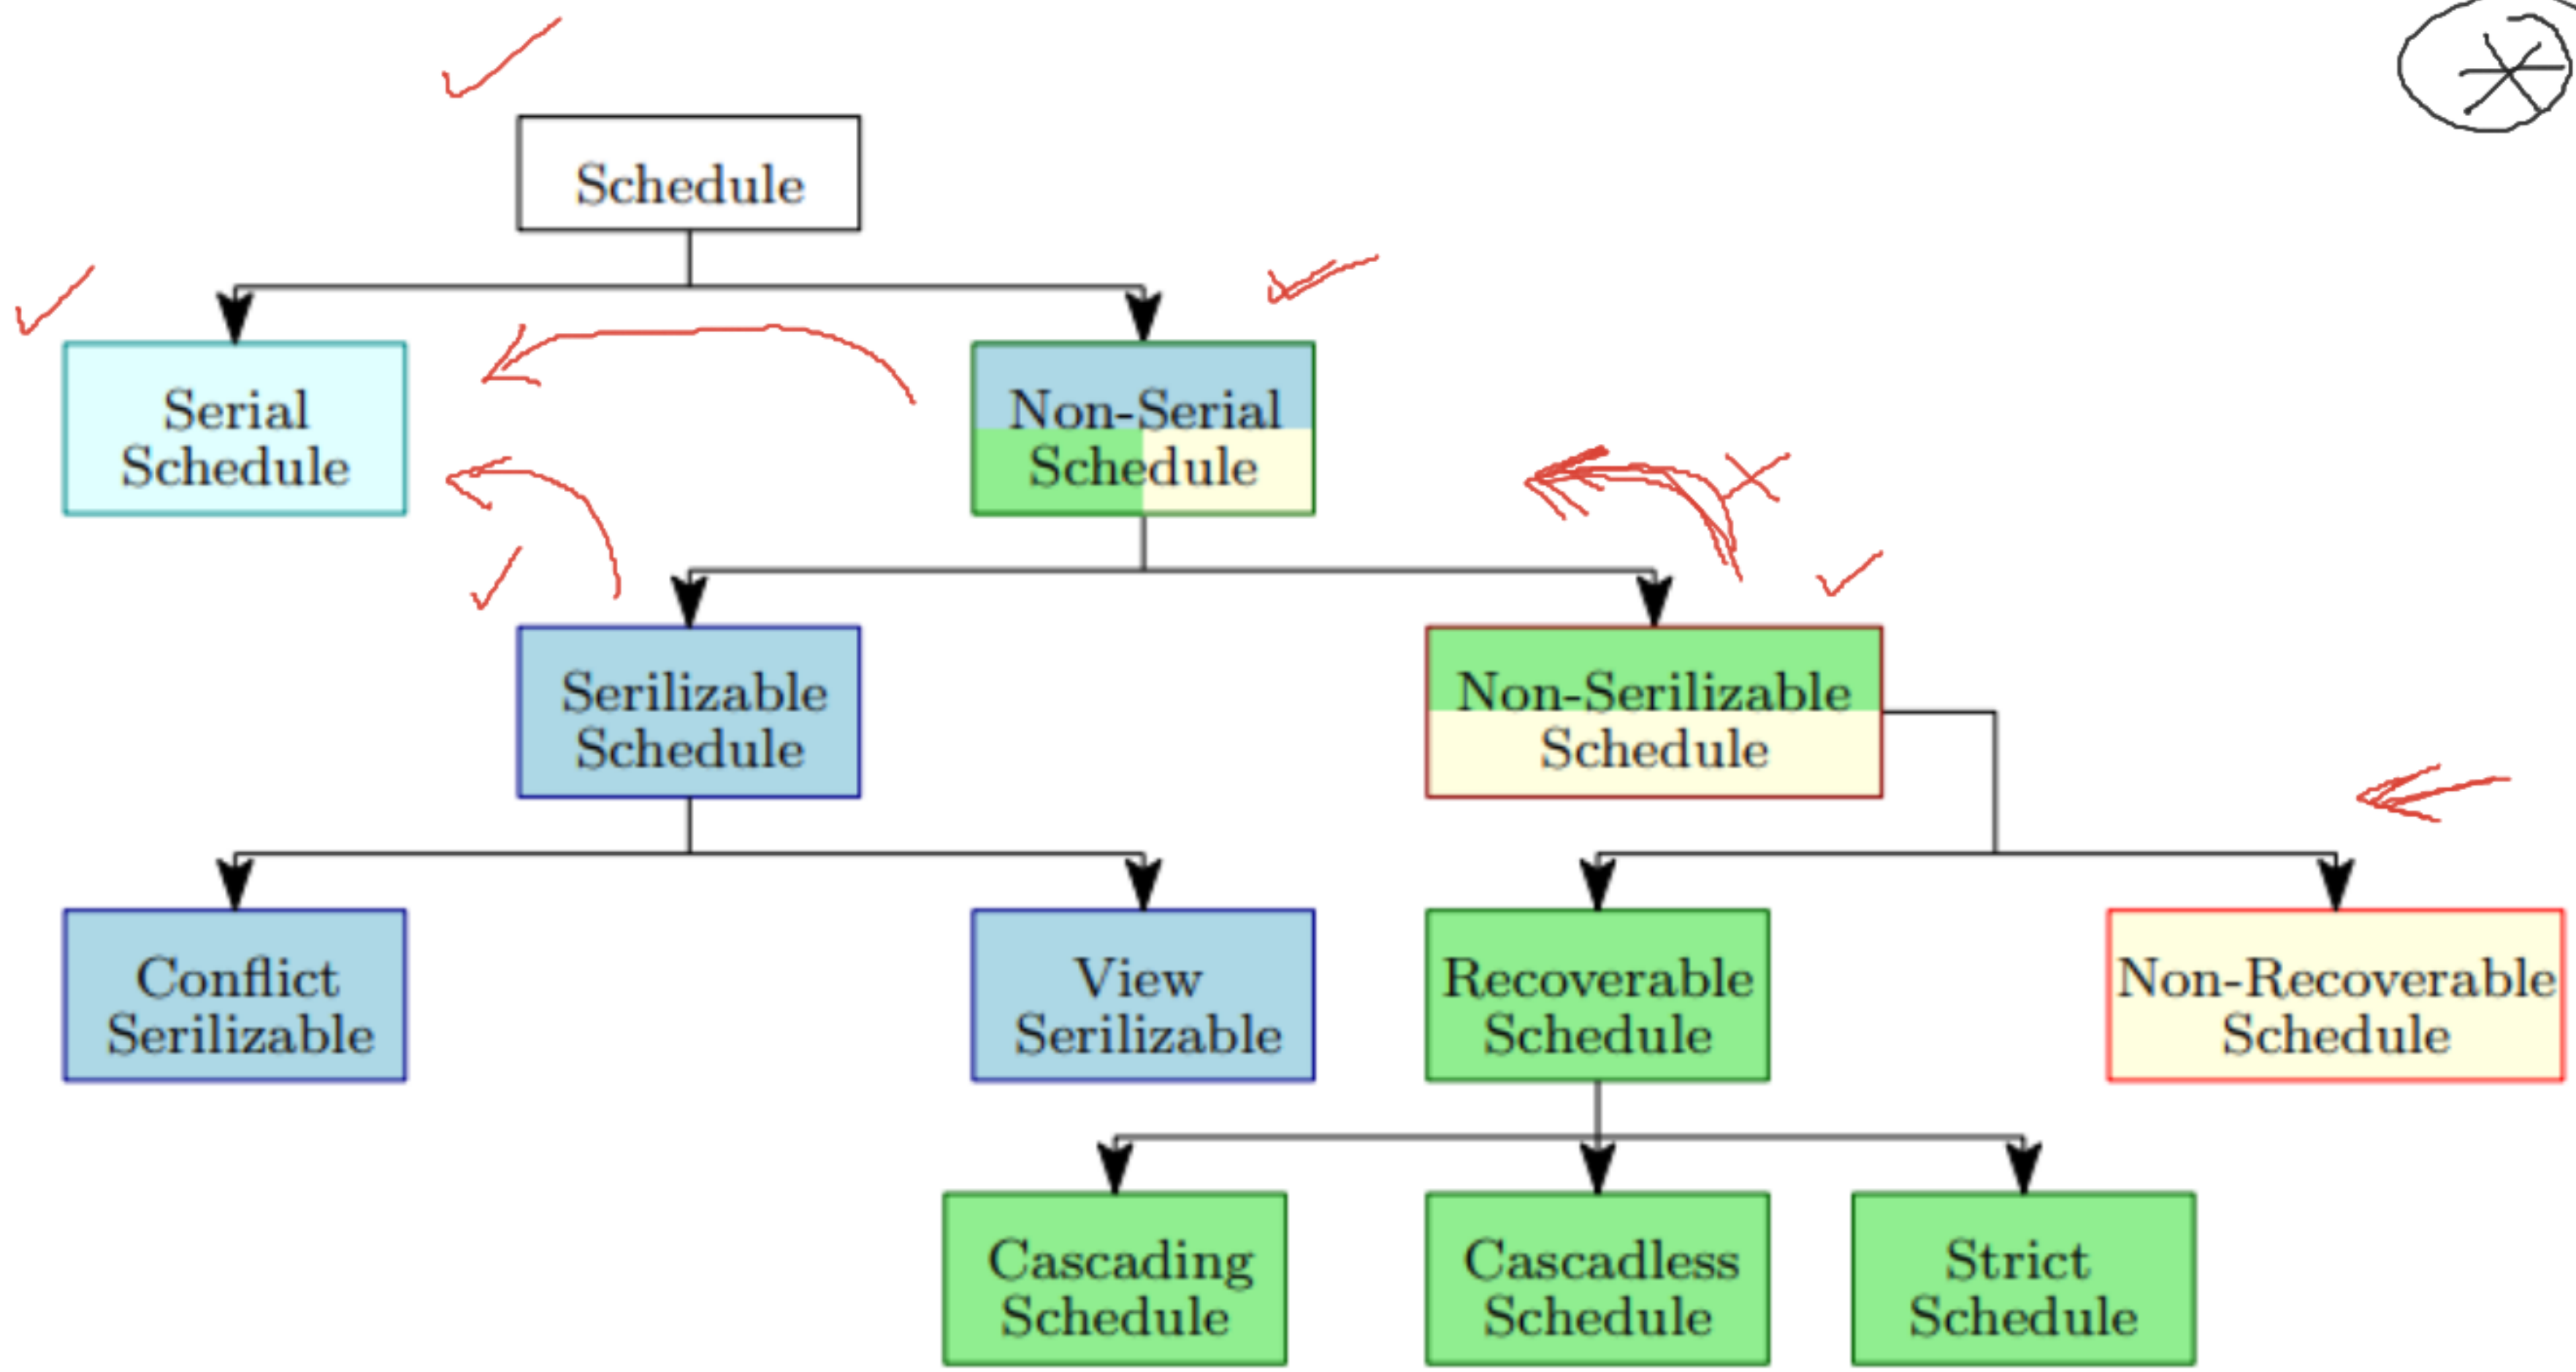
Suppose Initially,  $A = 1000$ .

After successful execution of  $T$ ,  $A = 500$

The changes of  $A$  should be permanent in the database

To ensure durability, multiple copies of database are stored at different locations





# Serial Schedule

- Transactions are executed in non-interleaved manner
- No transaction starts until a running transaction has ended

$T_1$	$T_2$
R(A) R(B) W(B) R(C) W(A)	
	R(A) R(B) W(A) R(C) W(C)

Figure: Example of a serial schedule

# Non-Serial Schedule

- Operations of multiple transactions are interleaved
- • The end result has to be correct and same as the serial schedule ← AIM
- • One transaction does not wait for another to complete
- Provide benefit of concurrent transaction
- • May lead to concurrency problem

	$T_1$	$T_2$
8 AM →	✓✓ R(A)	
8:01 →	✓ R(B)	
		✓ R(A) ← 8:02
	W(B)	
		R(B)
	R(C)	
		W(A)
		R(C)
		W(C)




Figure: Example of a non-serial schedule



## Serializable Schedule

- A non-serial schedule is said to be serializable only if it is equivalent to a serial schedule, for an n number of transactions
- It helps in improving resource utilization and CPU throughput

LHS		RHS	
Non-Serial Schedule		Equivalent Serial Schedule	
$T_1$	$T_2$	$T_1$	$T_2$
1 $\rightarrow R(A)$		1 $\rightarrow R(A)$	
2 $\rightarrow R(B)$		2 $\rightarrow R(B)$	
3 $\rightarrow W(B)$	3 $\rightarrow R(A)$	3 $\rightarrow W(B)$	
4 $\rightarrow R(C)$	4 $\rightarrow R(B)$	4 $\rightarrow R(C)$	
5 $\rightarrow W(A)$	5 $\rightarrow W(A)$		5 $\rightarrow R(A)$
6 $\rightarrow R(C)$	6 $\rightarrow R(C)$		6 $\rightarrow R(B)$
7 $\rightarrow W(C)$	7 $\rightarrow W(C)$		7 $\rightarrow W(A)$
8 $\rightarrow R(C)$	8 $\rightarrow W(C)$		8 $\rightarrow R(C)$
9 $\rightarrow W(C)$	9 $\rightarrow W(C)$		9 $\rightarrow W(C)$

Figure: Example of a serializable schedule

$$\begin{array}{l} \rightarrow L H J \quad S_1 \\ R_1(A) \\ R_1(B) \\ R_2(A) \\ W_1(B) \\ R_2(B) \\ R_1(C) \\ W_2(A) \\ R_2(C) W_2(C) \end{array}$$

# Conflict Seralizable Schedule

## Conflicting operations

Two operations are said to be conflicting if all the following conditions are satisfied:

- They belong to different transactions
- They operate on the same data item
- At Least one of them is a write operation

Example:

- ✗ •  $(R_1(A), W_2(A))$ : Conflicting operations. They belong to two different transactions on same data A and one of them is write operation
- ✗ → •  $(W_1(A), W_2(A))$ : Conflicting operations pair
- ✗ → •  $(W_1(A), R_2(A))$ : Conflicting operations pair
- ✓ → •  $(R_1(A), W_2(B))$ : Non-conflicting pair, as they operate on different data items
- ✓ •  $((W_1(A), W_2(B)))$ : Non-conflicting pair



# Conflict Seralizable Schedule (CSS)

## Conflict Serializable

A schedule is called conflict serializable if it can be transformed into a serial schedule by swapping non-conflicting operations

Example: Consider two following transactions

→  $T_1 : R_1(A), W_1(A), R_1(B), W_1(B)$  ←

→  $T_2 : R_2(A), W_2(A), R_2(B), W_2(B)$  ←

$S_1$  is CSS

Consider following schedule:

→  $S_1 : R_1(A), W_1(A), R_2(A), W_2(A), R_1(B), W_1(B), R_2(B), W_2(B)$

The schedule is conflict serializable

→  $S_{11} : R_1(A), W_1(A), R_2(A), W_2(A), R_1(B), W_1(B), R_2(B), W_2(B)$

→  $S_{12} : R_1(A), W_1(A), R_2(A), R_1(B), W_2(A), W_1(B), R_2(B), W_2(B)$

→  $S_{13} : R_1(A), W_1(A), R_1(B), R_2(A), W_2(A), W_1(B), R_2(B), W_2(B)$

$S_{14} : R_1(A), W_1(A), R_1(B), R_2(A), W_1(B), W_2(A), R_2(B), W_2(B)$

→  $S_{15} : R_1(A), W_1(A), R_1(B), W_1(B), R_2(A), W_2(A), R_2(B), W_2(B)$

$S_{15}$  is the equivalent serial schedule of  $S_1$



Operations

$R(A)$   $W(A)$

Transactions

$T_1, T_2$

Schedule

$S_1, S_2$

## Conflict Equivalent

### Conflict Equivalent

Two schedules are said to be conflict equivalent when one can be transformed to another by swapping non-conflicting operations

Example: Consider two following transactions

$T_1 : R_1(A), W_1(A), R_1(B), W_1(B)$

$T_2 : R_2(A), W_2(A), R_2(B), W_2(B)$

Consider following two schedules:

$S_1 : R_1(A), W_1(A), R_2(A), W_2(A), R_1(B), W_1(B), R_2(B), W_2(B)$

$S_2 : R_1(A), W_1(A), R_1(B), W_1(B), R_2(A), W_2(A), R_2(B), W_2(B)$

$S_1$  and  $S_2$  are conflict equivalent